

The Entity-Relationship Model

Chapter 2



Databases Model the Real World

■ Data Model:

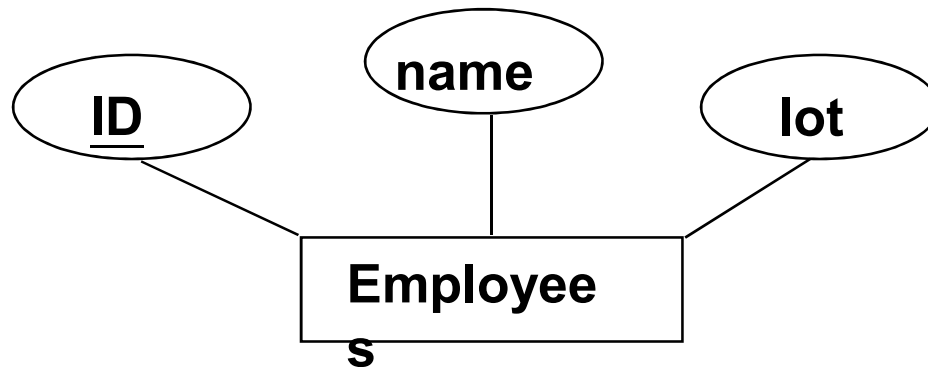
A *data Model* is a collection of conceptual tools for describing data, data relationships, data semantics and consistency constraints.

■ E-R Model:

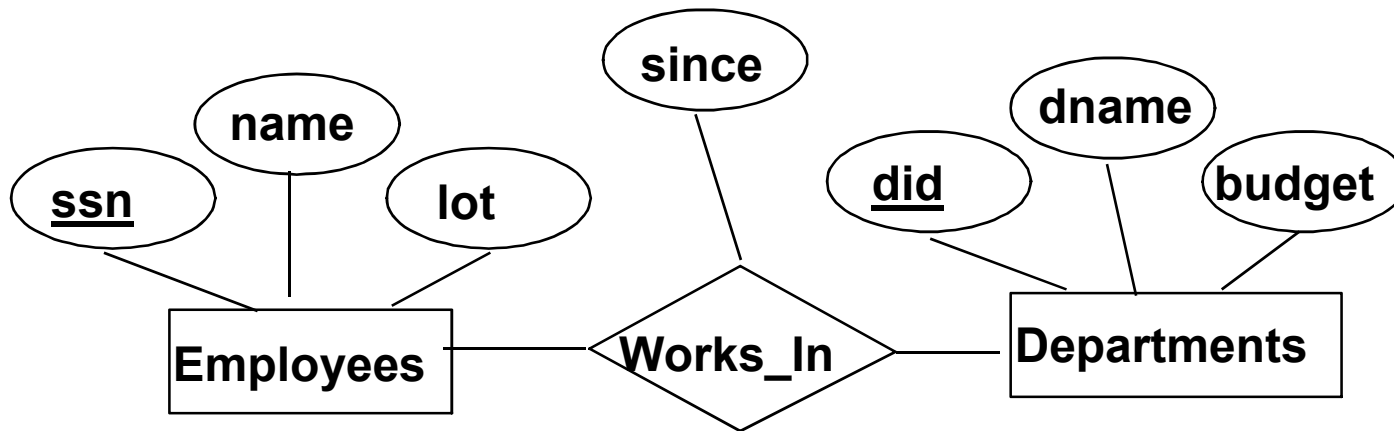
The Entity-Relationship data model is based on a perception of a real world that consists of a set of basic objects called entities and of relationships among these objects.

E-R Model

- **Entity:** An entity is an object that exists in the real world and is distinguishable from other objects. An entity is described using a set of attributes.
- **Attribute:** Attributes are descriptive properties possessed by each member of an entity set.
- **Entity Set:** A collection of all entities of the same type is an entity set. E.g. all employees.
 - All entities in an entity set have the same set of attributes.



ER Model Basics



- **Relationship**: A relationship is an association among two or more entities. relationships can have their own attributes.
Relationships indicate a meaningful connection between two entity types.
- **Relationship Set**: The Collection of all similar relationships of the same type is called relationship set.

Attribute & Constraint

- Attributes are facts, aspects, properties, or details about an entity
 - Students have IDs, names, courses, addresses, ...
 - Modules have codes, titles, credit weights, levels, ...

Constraint:

Constraints are restrictions on legal relation states

Types of Attribute

- **Simple-value attribute:**

These attributes are not divided into subparts.

- **Composite attributes:**

They can be divided into subparts.

Example: name which contain first-name, middle-name and last-name.

- **Single-valued attribute:**

The attribute which specify only one number is called single-valued attribute.

Exam: loan-number which indicate only one loan number.

- **Multi-valued attribute:**

When an attribute has a set of values for specific entity is called multi-valued attribute.

Exam: Phone-number- An employee may have one or more phone numbers.

- **Derived Attribute:** The value for this type of attribute can be derived from the values of other related attributes or entities.

Example: *Date-of birth* --- we can calculate age from ***date-of-birth*** and the current date. Thus age is a derived attribute.

Loans_held, which represents how many loans a customer has from the bank.

We can derive the value for this attribute by counting the number of loan entities associated with that customer.

- **Null value:** An attribute takes a null value when an entity does not have a value for it, that means the value does not exist for the entity.

Example: *middle_name* – One may have no middle name.

Mapping Cardinalities

- **Mapping cardinalities, or cardinality ratios, express the number of entities to which another entity can be associated via a relationship set.**
 - Mapping cardinalities are most useful in describing binary relationship sets.
- **There are four types of binary relationship in E-R model:**
 - ***One to many :***

An entity in A is associated with at most one entity in B and an entity in B is associated with at most one entity in A.
 - ***One to Many:***

An entity in A is associated with any number of (zero or more) entities in B. An entity in B can be associated with at most one entity in A.

Types of E-R model

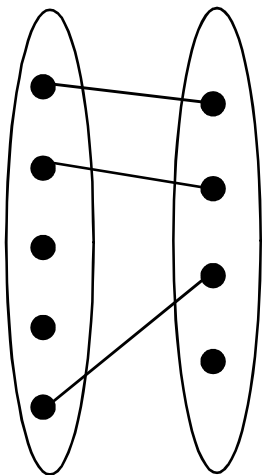
- **Many to One :**

An entity in A is associated with at most one entity in B. An entity in B, however, can be associated with any number (zero or more) of entities in A.

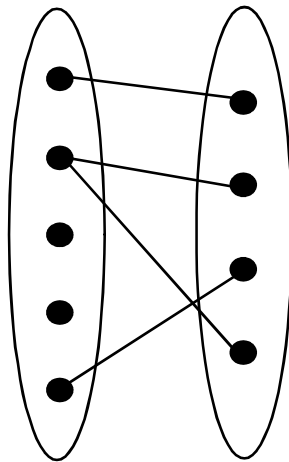
- **Many to Many :**

An entity in A is associated with any number (zero or more) of entities in B and an entity in B is associated with any number (zero or more) of entities in A.

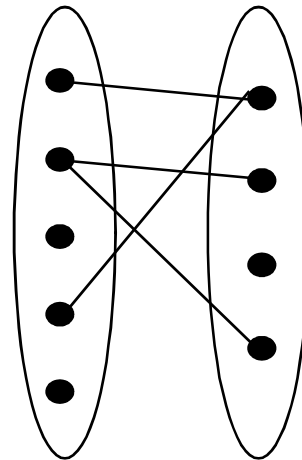
Types of E-R model



One to one



one-to Many



Many-to-Many

Keys

- **The key allow us to identify a set of attributes that suffice to distinguish entities from each other.**
- Keys are used to identify particular entities.
- Keys also help uniquely identify relationships and this distinguish relationships from each other.
- **There are many types of keys:**
 - ❖ **Superkey**
 - ❖ **Candidate key**
 - ❖ **Primary key**
 - ❖ **Foreign key**


Superkey

- ❖ **A superkey is a set of one or more attributes that, taken collectively allow us to identify uniquely an entity in the entire set.**
 - Example: customer_id attribute of the entity set customer is sufficient to distinguish one customer entity to another.
 - The combination of customer_name and customer_id is a superkey for the entity set customer.
 - The customer_name attribute of customer is not a superkey, because several people might have the same name.

■ Candidate Key:

- ✓ A candidate is a subset of a super key. A candidate key is a single field or the least combination of fields that uniquely identifies each record in the table.
- ✓ Every table must have at least one candidate key but at the same time can have several.

Candidate Keys



StudentId	firstName	lastName	courseId
L0002345	Jim	Black	C002
L0001254	James	Harradine	A004
L0002349	Amanda	Holland	C002
L0001198	Simon	McCloud	S042
L0023487	Peter	Murray	P301
L0018453	Anne	Norris	S042

Primary key

- ❖ It is a candidate key that is chosen by the database designer to identify entities within an entity set. Primary key is a single attribute or combination of attributes that uniquely defines a database record
- The primary key should be chosen such that its attributes are never, or very rarely, changed.

Primary Keys



<u>StudentId</u>	firstName	lastName	courseId
L0002345	Jim	Black	C002
L0001254	James	Harradine	A004
L0002349	Amanda	Holland	C002
L0001198	Simon	McCloud	S042
L0023487	Peter	Murray	P301
L0018453	Anne	Norris	S042

Foreign Key

- ❖ **A foreign key is a field (or collection of fields) in one table that uniquely identifies a row of another table.**
- A foreign key is generally a primary key from one table that appears as a field in another where the first table has a relationship to the second.

Example: if we had a table A with a primary key X that linked to a table B where X was a field in B, then X would be a foreign key in B.

- **The table containing the foreign key is called the referencing or child table.**
- **The table containing the candidate key is called the referenced or parent table.**

Foreign key

<u>studentId</u>	firstName	lastName	courseId
L0002345	Jim	Black	C002
L0001254	James	Harradine	A004
L0002349	Amanda	Holland	C002
L0001198	Simon	McCloud	S042

Foreign Keys

Relationship

Primary Keys

<u>courseId</u>	courseName
A004	Accounts
C002	Computing
P301	History
S042	Short Course



- **Secondary key:** All the rest of candidate keys .

- **Composite Key:**

Composite key consists of more than one attributes.

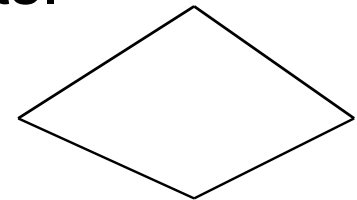
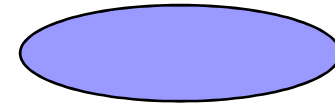
- **Tuple:**

In the context of databases, a tuple is one record (one row)

In other word, Tuples are *unordered* sets of known values with names.

E-R diagram

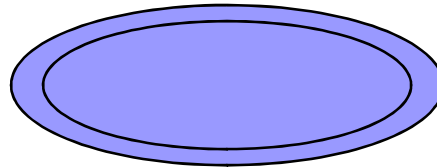
- E-R diagram can express the overall logical structure of a database graphically. E-R diagrams are simple and clear.
- ✓ Rectangles – which represent entity sets.
- ✓ Ellipses – which represent attributes.
- ✓ Diamond – which represent relationship sets.



- ✓ Lines-which link attributes to entity sets and entity sets to relationship sets.



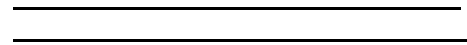
- ✓ Double ellipses- which represent multivalued attributes.



- ✓ Dashed ellipses-which denote derived attributes.

E-R Diagram

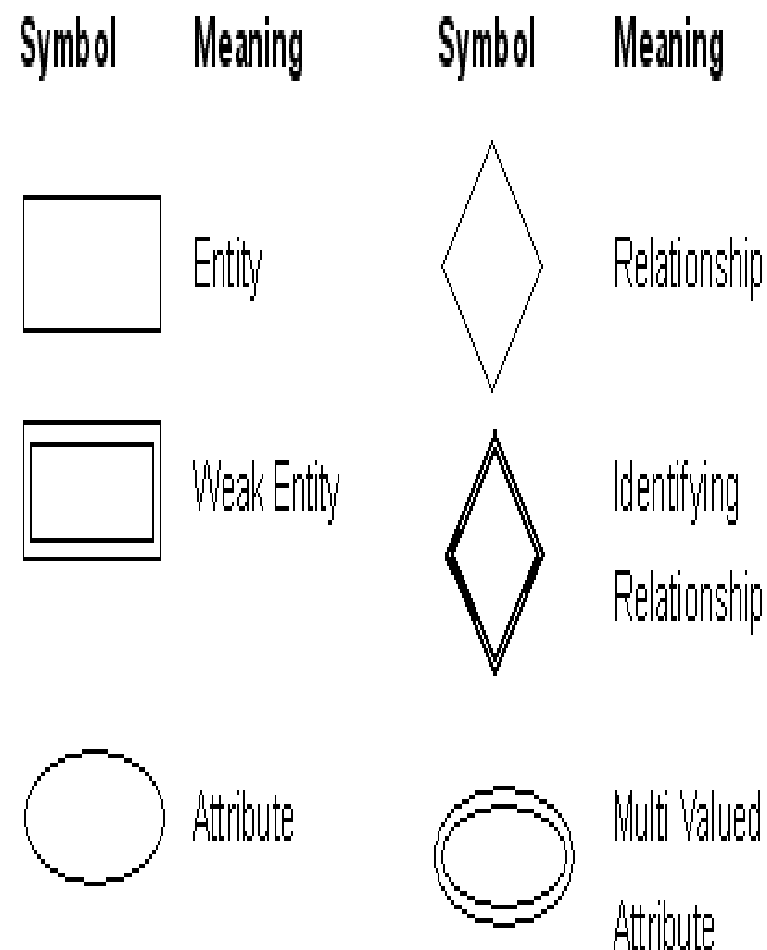
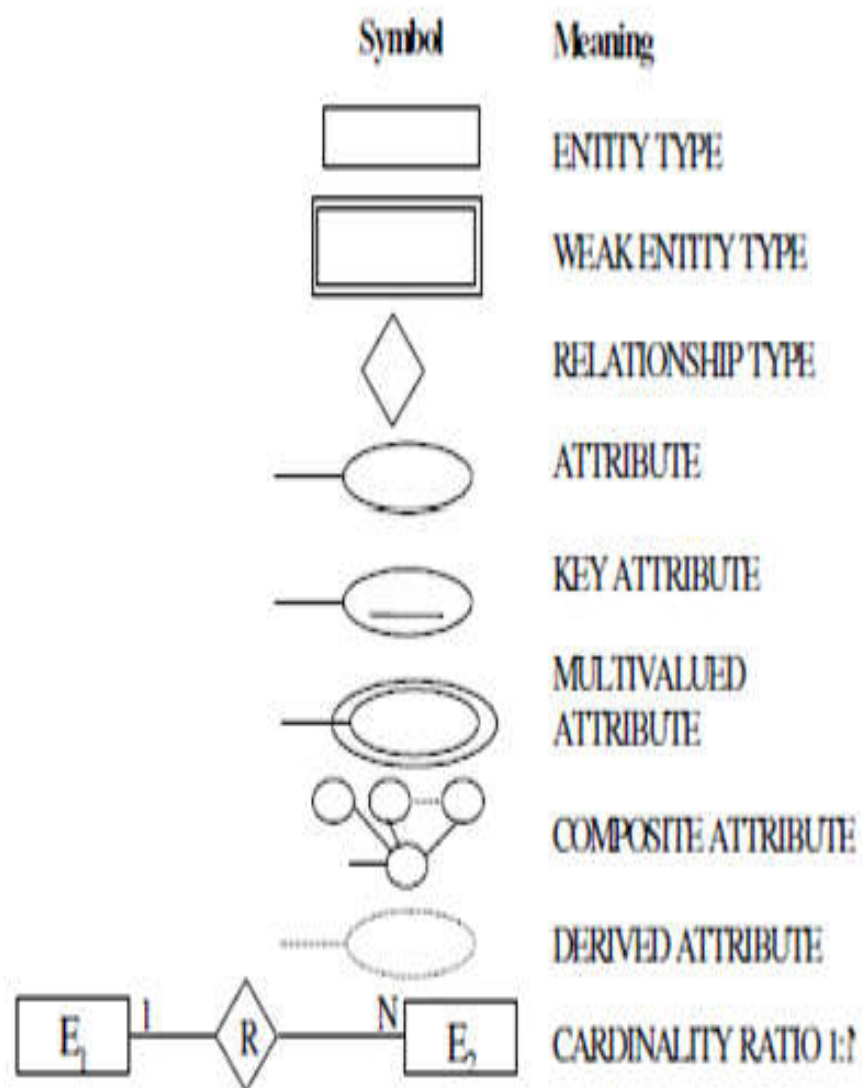
- ✓ Double lines-which indicate total participation of an entity in a relationship set.



- ✓ Double rectangle:-which represent weak entity sets .

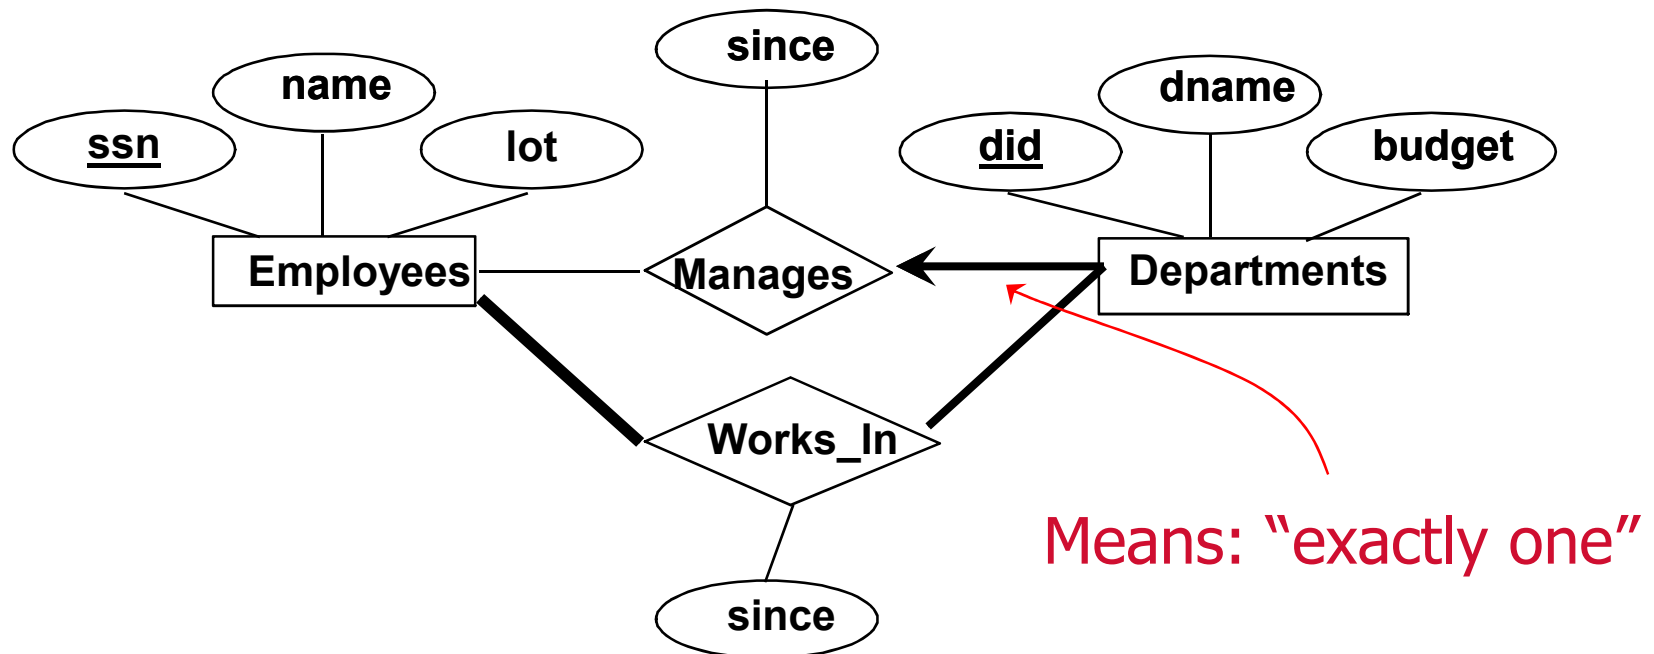


E-R Diagram



Participation Constraints

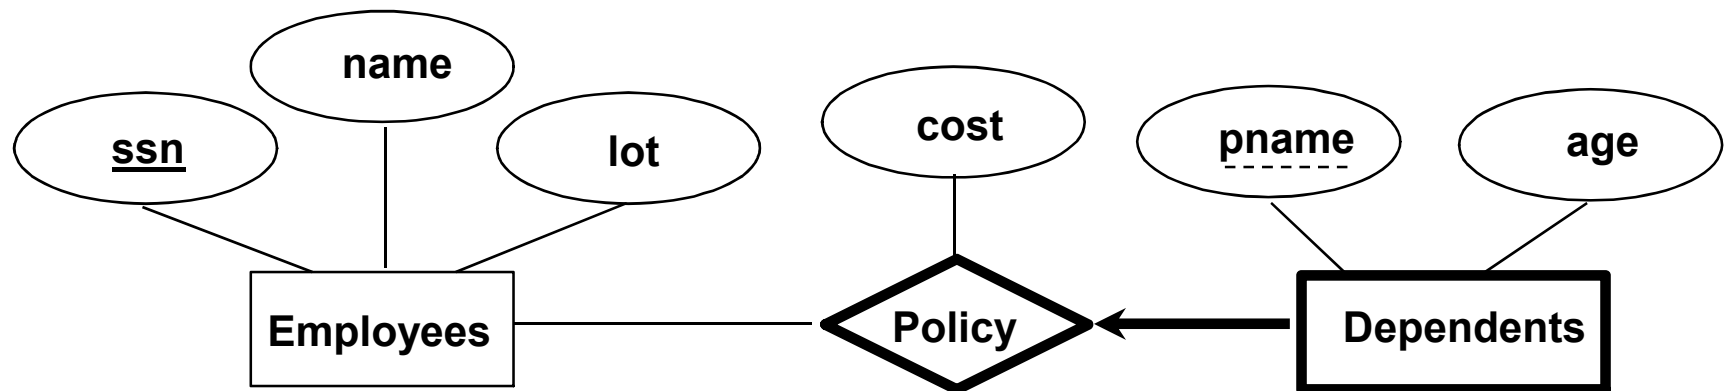
- Does every employee work in a department?
- If so, this is a *participation constraint*
 - the participation of Employees in Works_In is said to be *total* (vs. *partial*)
 - What if every department has an employee working in it?
- Basically means “at least one”



Weak Entities

A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.

- Owner entity set and weak entity set must participate in a one-to-many relationship set (one owner, many weak entities).
- Weak entity set must have total participation in this *identifying* relationship set.

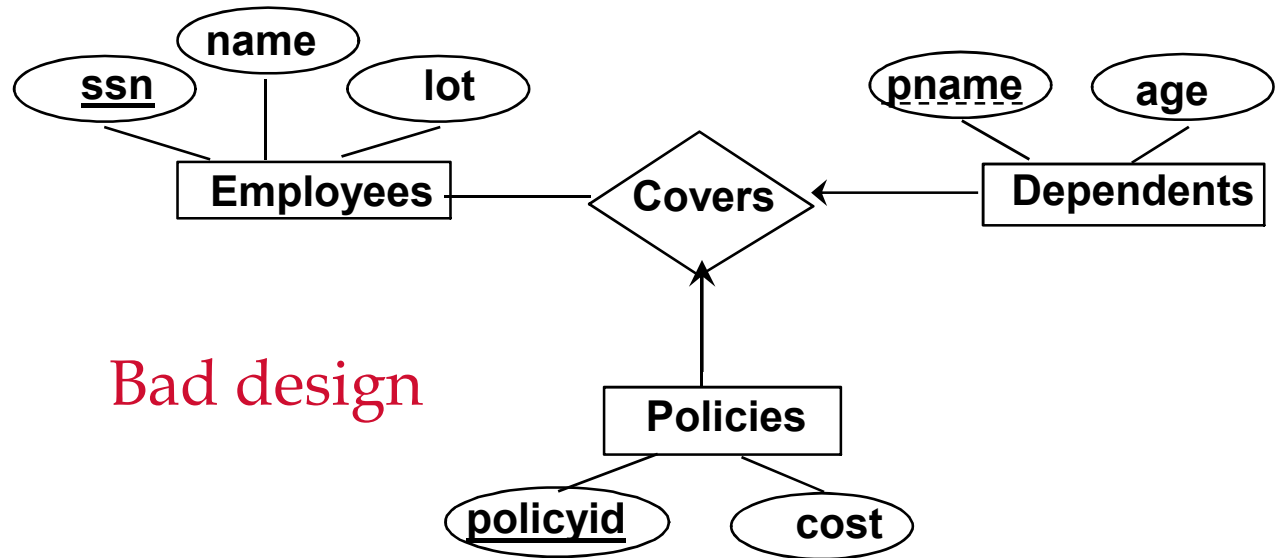


Weak entities have only a “partial key” (dashed underline)

Binary vs. Ternary Relationships

If each policy is owned by just 1 employee:

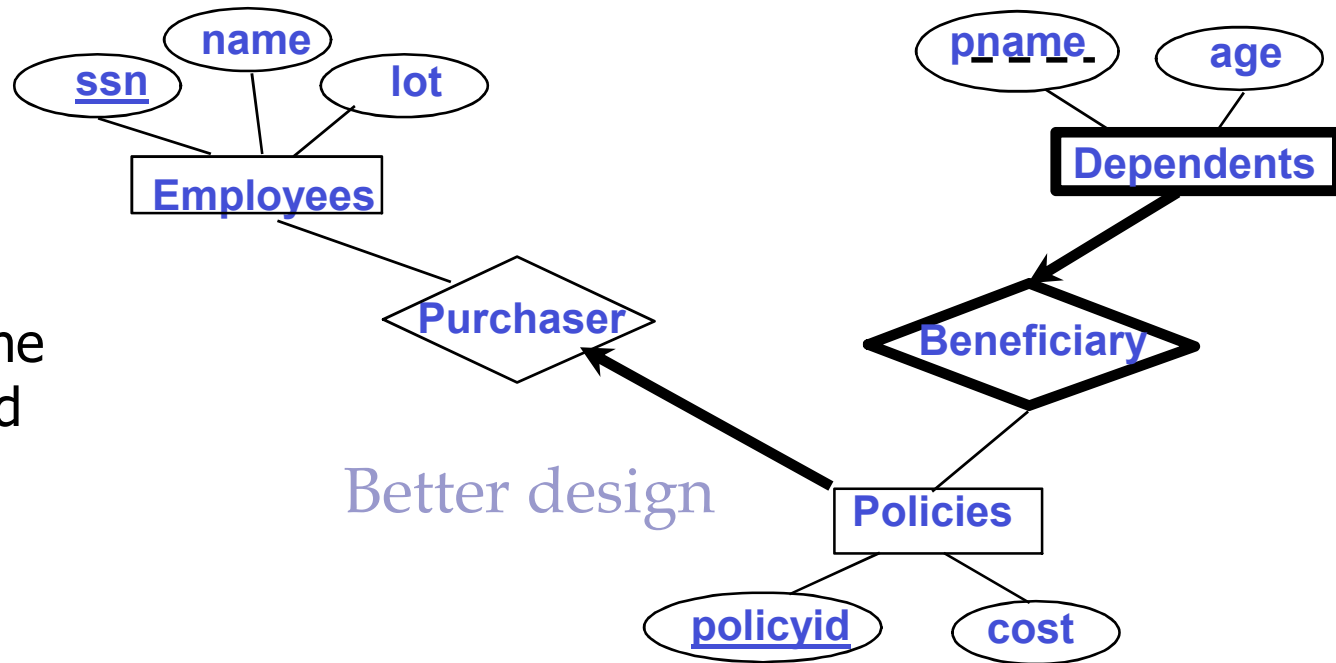
Bad design



Key constraint on Policies would mean policy can only cover 1 dependent!

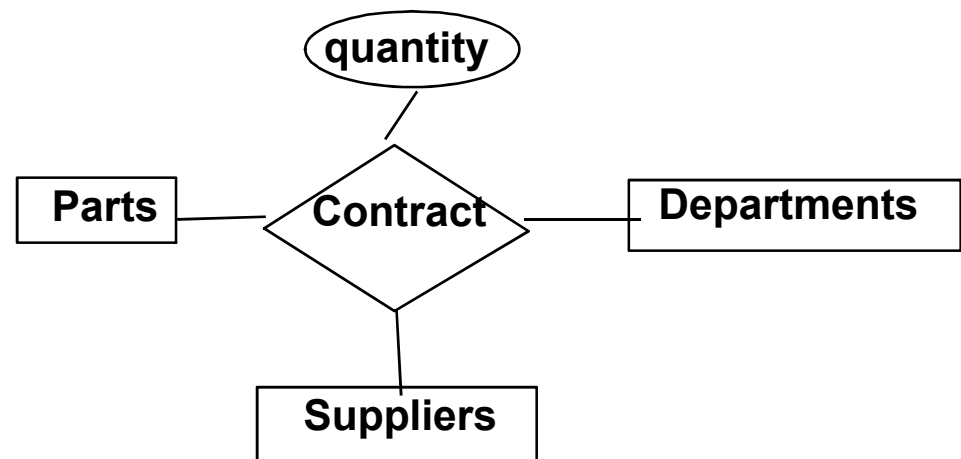
- Think through *all* the constraints in the 2nd diagram!

Better design

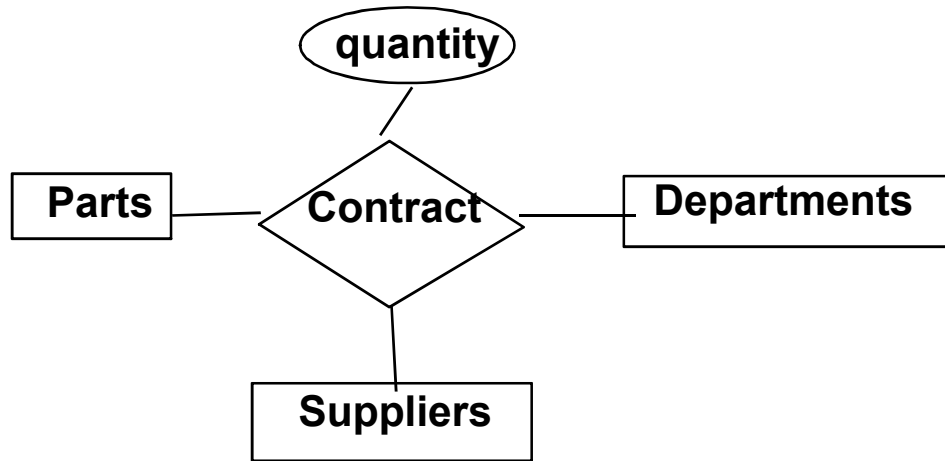


Binary vs. Ternary Relationships

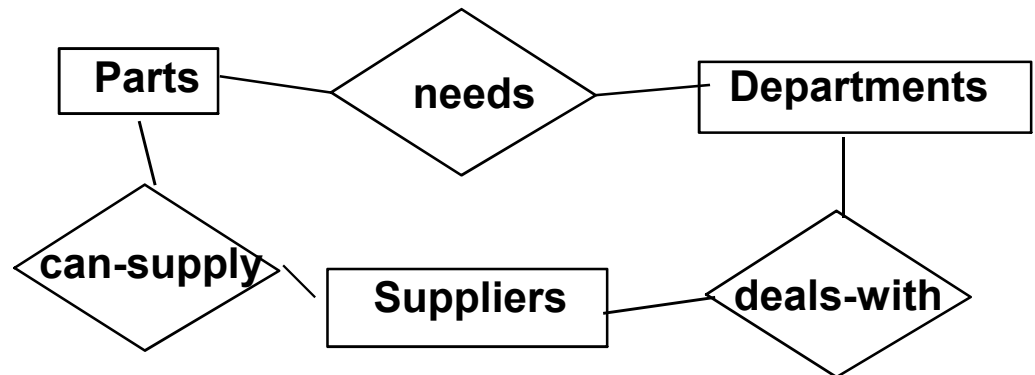
- Previous example illustrated a case when two binary relationships were better than one ternary.
- An example in the other direction: a ternary relation **Contracts** relates entity sets **Parts**, **Departments** and **Suppliers**, **and has descriptive attribute *quantity***.
 - No combination of binary relationships is an adequate substitute.



Binary vs. Ternary Relationships



VS.

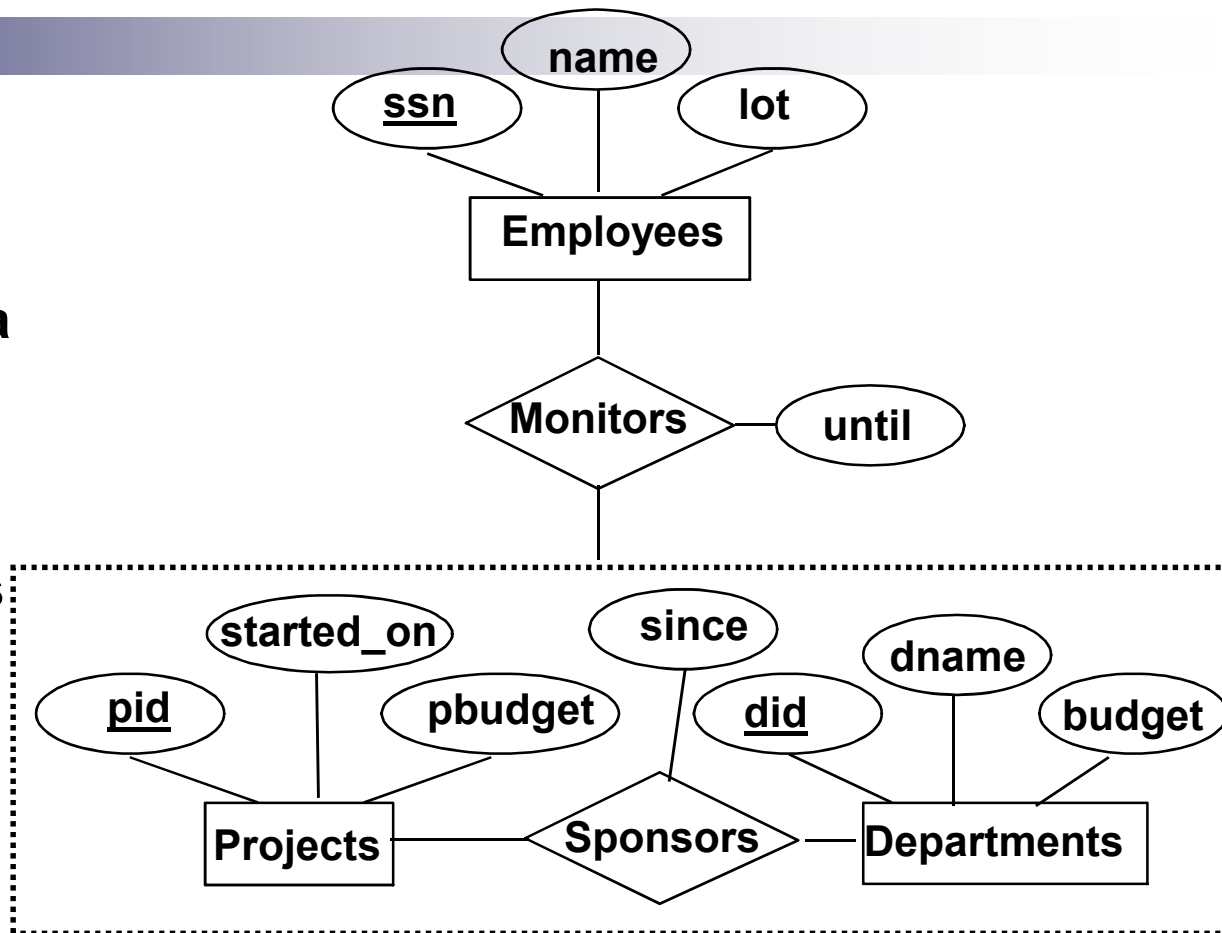


- S “can-supply” P, D “needs” P, and D “deals-with” S does not imply that D has agreed to buy P from S.
- How do we record *qty*?

Aggregation

Used to model a relationship involving a *relationship set*.

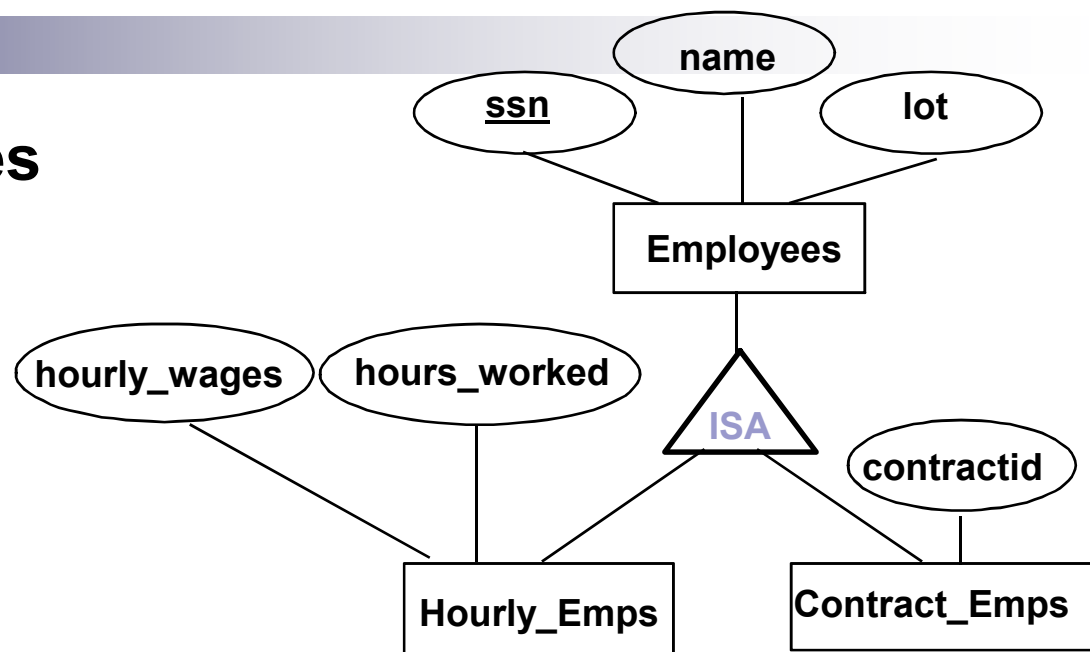
Allows us to **treat a relationship set as an entity set** for purposes of participation in (other) relationships.



Aggregation vs. ternary relationship?

- ❖ Monitors is a distinct relationship, with a descriptive attribute.
- ❖ Also, can say that each sponsorship is monitored by at most one employee.

ISA ('is a') Hierarchies



- *Overlap constraints*: Can Simon be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
- *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)
- Reasons for using ISA:
 - To add descriptive attributes specific to a subclass.
 - i.e. not appropriate for all entities in the superclass
 - To identify entities that participate in a particular relationship
 - i.e., not all superclass entities participate

Review - Our Basic ER Model

- Entities and Entity Set (boxes)
- Relationships and Relationship sets (diamonds)
 - binary
 - n-ary
- Key constraints (1-1, 1-M, M-M, arrows on 1 side)
- Participation constraints (bold for Total)
- Weak entities - require strong entity for key
- Aggregation - an alternative to n-ary relationships
- Isa hierarchies - abstraction and inheritance

Conceptual Design Using the ER Model

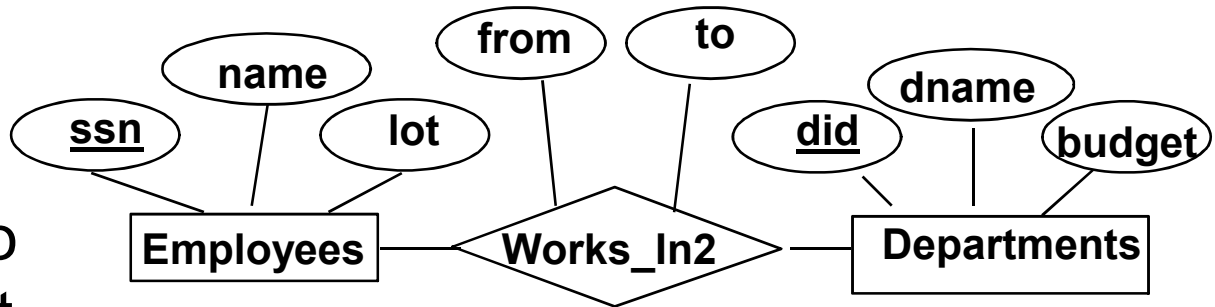
- ER modeling *can* get tricky!
- Design choices:
 - Should a concept be modeled as an **entity or an attribute**?
 - Should a concept be modeled as an **entity or a relationship**?
 - Identifying relationships: **Binary or ternary**? **Aggregation**?
- Note constraints of the ER Model:
 - A lot of data semantics can (and should) be captured.
 - But some constraints cannot be captured in ER diagrams.
 - We'll refine things in our logical (relational) design

Entity vs. Attribute

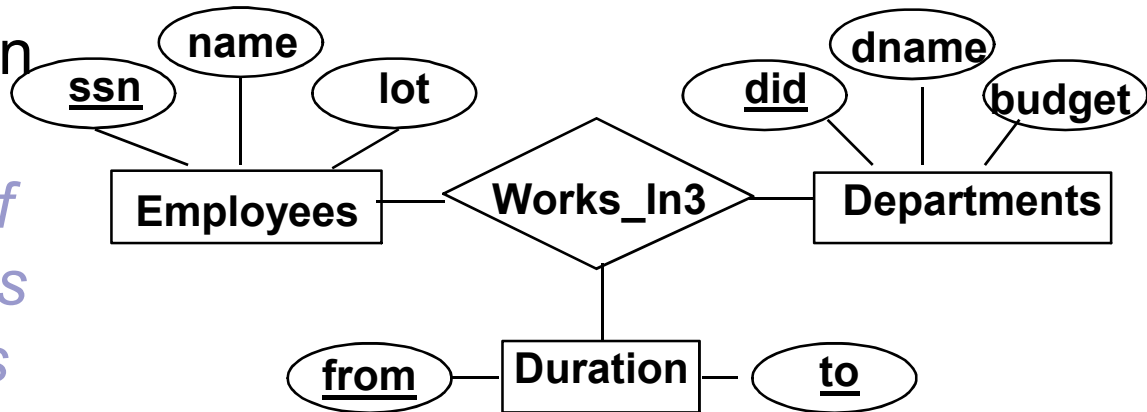
- Should *address* be an attribute of Employees or an entity (related to Employees)?
- **Depends** upon how we want to use address information, and the semantics of the data:
 - If we have **several addresses per employee**, *address* must be an entity (since attributes cannot be set-valued).
 - If the **structure** (city, street, etc.) **is important**, *address* must be modeled as an entity (since attribute values are atomic).

Entity vs. Attribute (Cont.)

- Works_In2 does not allow an employee to work in a department for two or more periods.

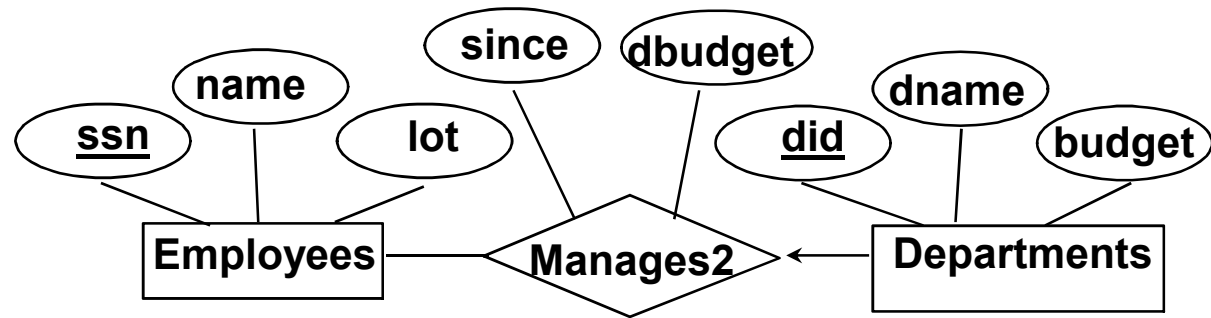


- Similar to the problem of wanting to record several addresses for an employee: we want to record *several values of the descriptive attributes for each instance of this relationship*.

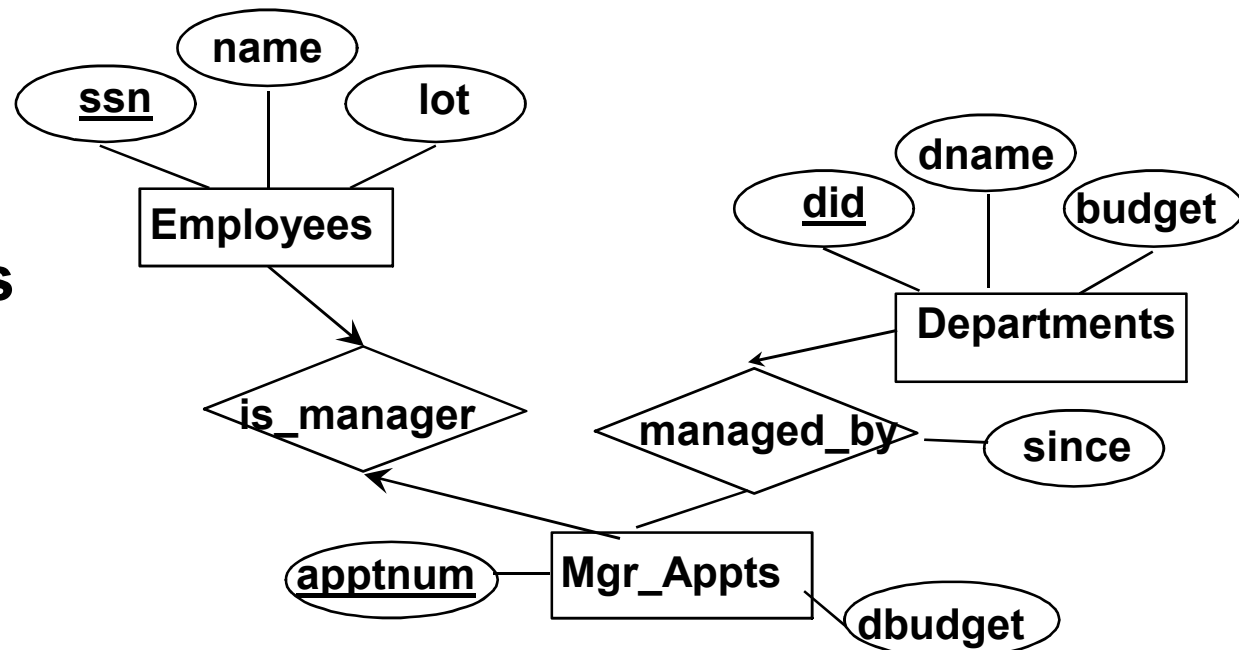


Entity vs. Relationship

OK as long as a manager gets a separate discretionary budget (*dbudget*) for each dept.



What if manager's *dbudget* covers *all* managed depts?
(can repeat value, but such redundancy is problematic)

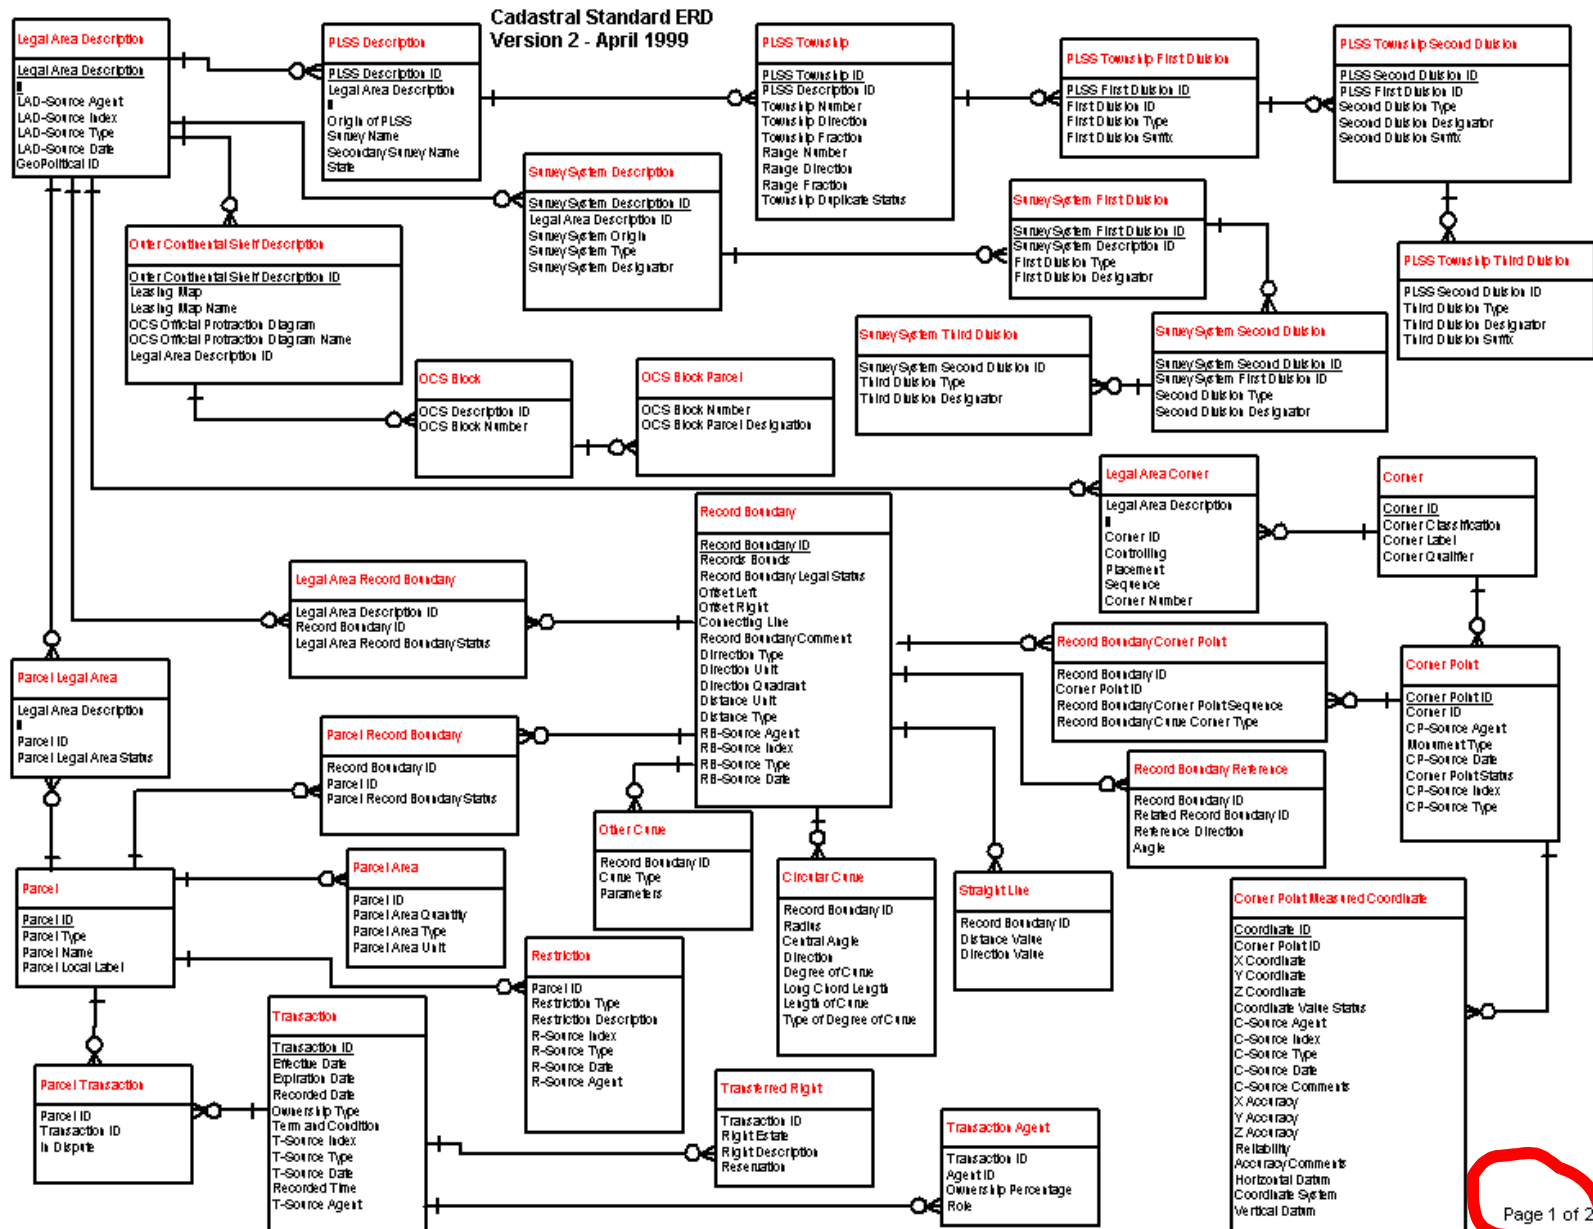




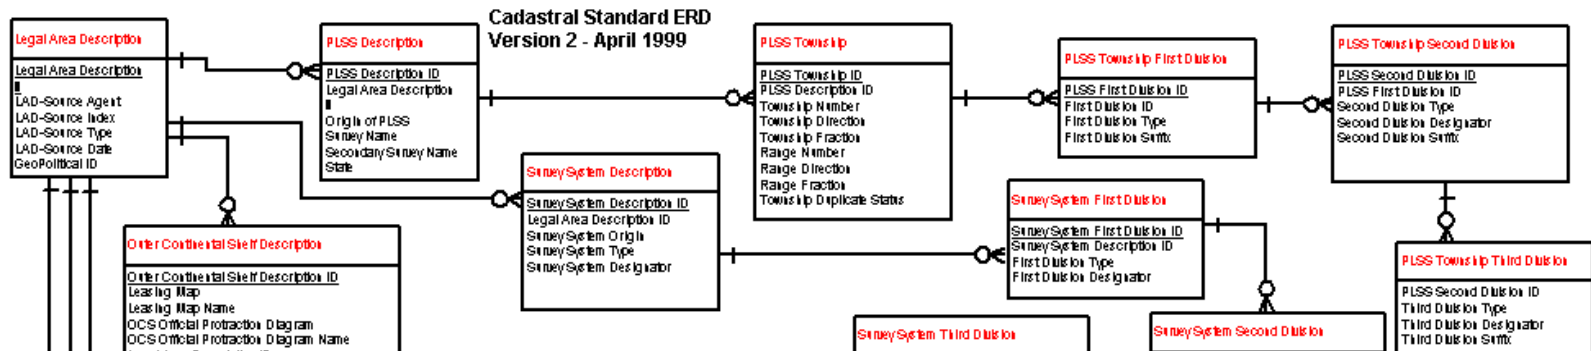
These things get pretty hairy!

- Many E-R diagrams cover entire walls!
- A modest example:

A Cadastral E-R Diagram

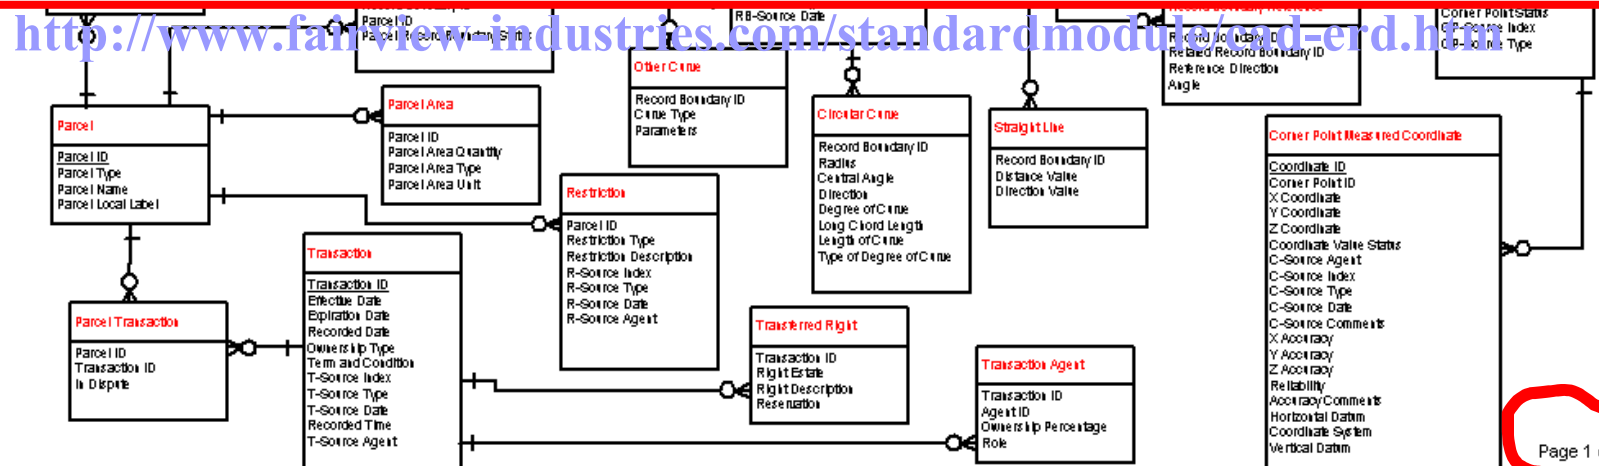


A Cadastral E-R Diagram



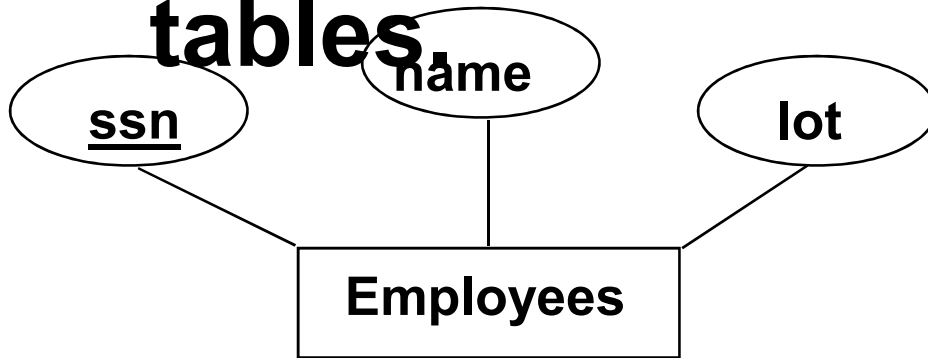
cadastral: showing or recording property boundaries, subdivision lines, buildings, and related details

Source: US Dept. Interior Bureau of Land Management, Federal Geographic Data Committee Cadastral Subcommittee



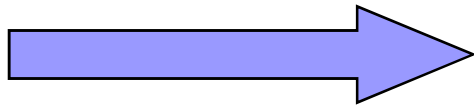
Logical DB Design: ER to Relational

■ Entity sets to tables



ssn	name	lot
123-22-3666	Attishoo	48
231-31-5368	Smiley	22
131-24-3650	Smethurst	35

```
CREATE TABLE Employees
(ssn CHAR(11),
 name CHAR(20),
 lot INTEGER,
PRIMARY KEY (ssn))
```



Relationship Sets to Tables

■ In translating a **many-to-many** relationship set to a relation, attributes of the relation must include:

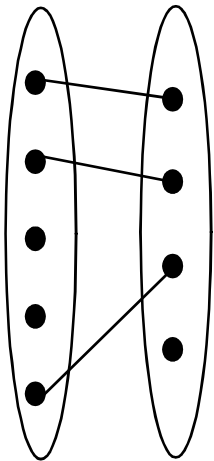
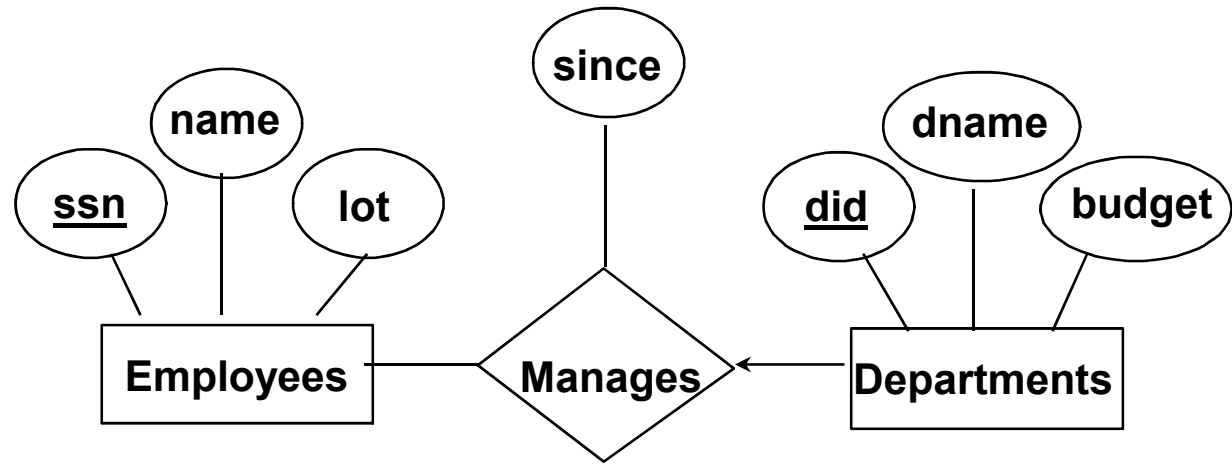
- 1) Keys for each participating entity set (as foreign keys). This set of attributes forms a *superkey* for the relation.
- 2) All descriptive attributes.

```
CREATE TABLE Works_In(  
    ssn CHAR(1),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees,  
    FOREIGN KEY (did)  
        REFERENCES Departments)
```

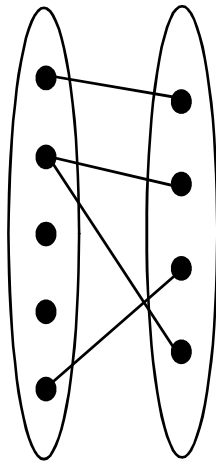
ssn	did	since
123-22-3666	51	1/1/91
123-22-3666	56	3/3/93
231-31-5368	51	2/2/92

Review: Key Constraints

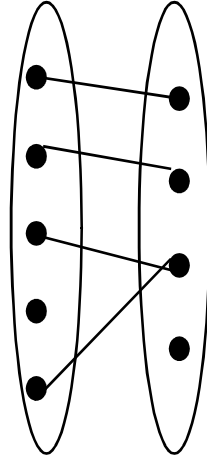
- Each dept has at most one manager, according to the key constraint on Manages.



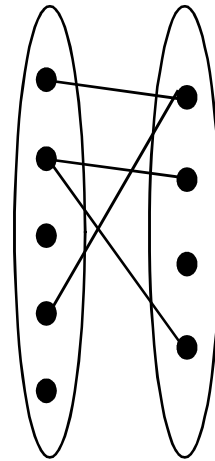
1-to-1



1-to Many



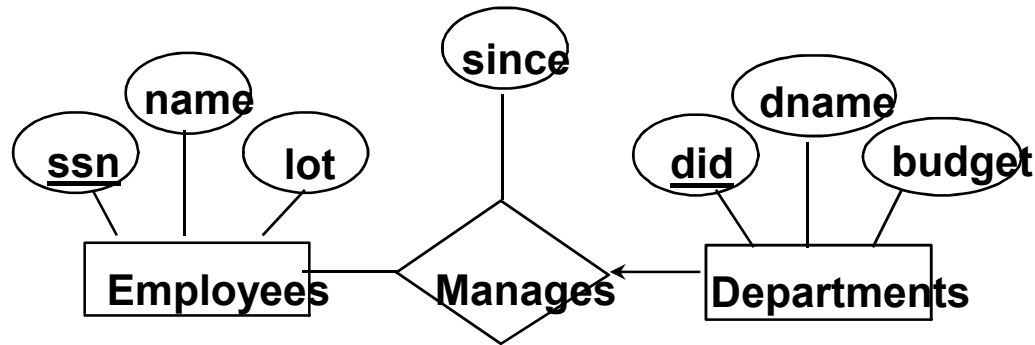
Many-to-1



Many-to-Many

Translation to relational model?

Translating ER with Key Constraints



- Since each department has a unique manager, we could instead combine **Manages** and **Departments**.

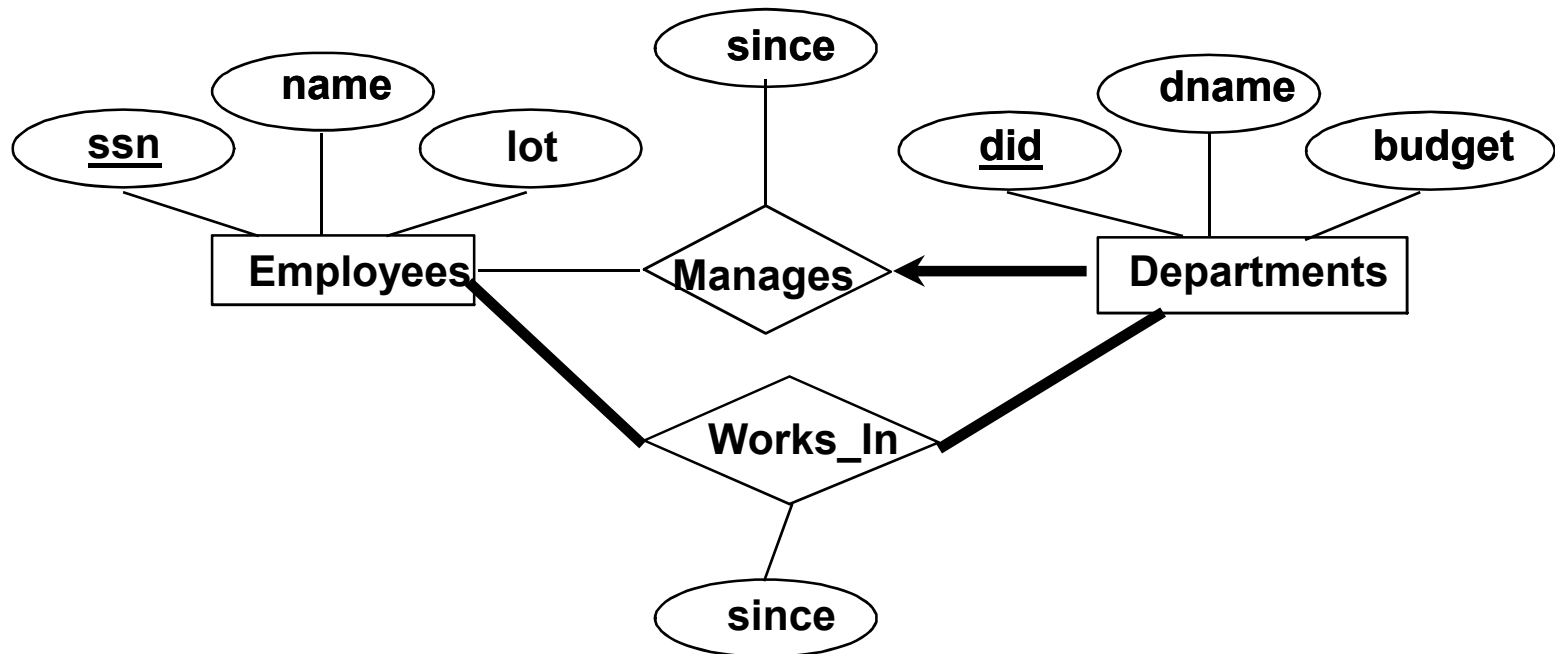
```
CREATE TABLE Manages(  
  ssn CHAR(11),  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn)  
  REFERENCES Employees,  
  FOREIGN KEY (did)  
  REFERENCES Departments)
```

Vs.

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11),  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn)  
  REFERENCES Employees)
```


Review: Participation Constraints

- Does every department have a manager?
 - If so, this is a participation constraint: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value!)



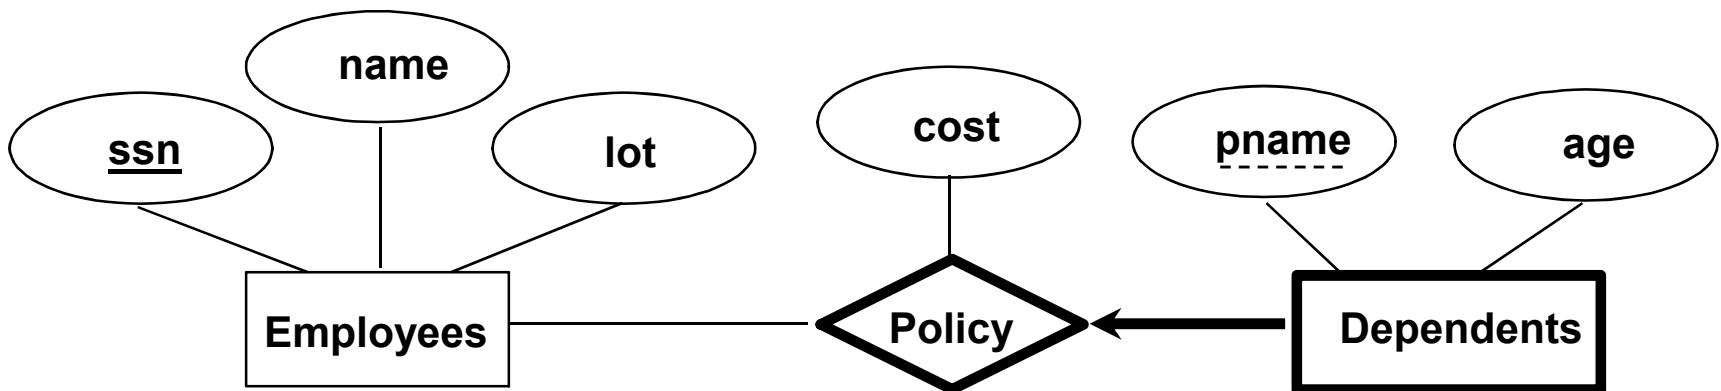
Participation Constraints in SQL

- We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```

Review: Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.

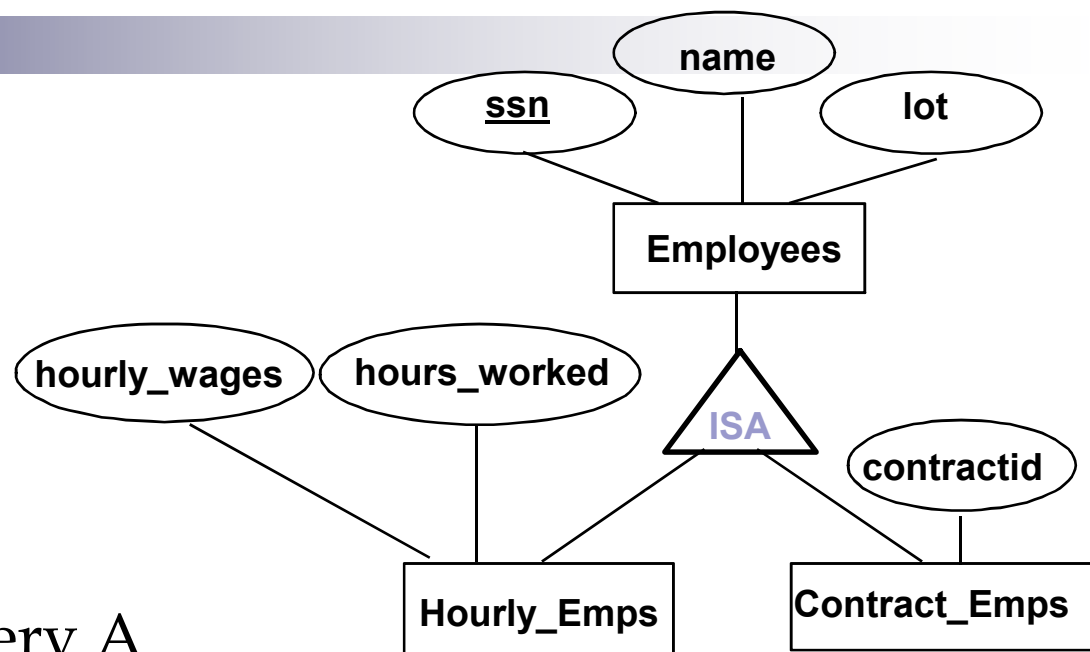


Translating Weak Entity Sets

- Weak entity set and identifying relationship set are translated into a single table.
 - When the owner entity is deleted, all owned weak entities must also be deleted.

```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

Review: ISA Hierarchies



❖ As in C++, or other PLs, attributes are inherited.

❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.

- *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
- *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)

Translating ISA Hierarchies to Relations

■ ***General approach:***

- 3 relations: Employees, Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*: Every employee is recorded in Employees. For hourly emps, extra info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*); must delete Hourly_Emps tuple if referenced Employees tuple is deleted).
 - Queries involving all employees easy, those involving just Hourly_Emps require a join to get some attributes.
- Alternative: Just Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
 - Each employee must be in one of these two subclasses.

Now you try it

University database:

- Courses, Students, Teachers
- Courses have ids, titles, credits, ...
- Courses have multiple sections that have time/rm and exactly one teacher
- Must track students' course schedules and transcripts including grades, semester taken, etc.
- Must track which classes a professor has taught
- Database should work over multiple semesters

Other SQL DDL Facilities

- Integrity Constraints (ICs) - Review
- An IC describes conditions that every *legal instance* of a relation must satisfy.
 - Inserts/deletes/updates that violate IC's are disallowed.
 - Can be used to ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be < 200)
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.
 - *Domain constraints*: Field values must be of right type. Always enforced.
 - *Primary key and foreign key constraints*: you know them.

General Constraints

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 ))
```

CREATE TABLE Reserves

```
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK (`Interlake' <>
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid))))
```

Constraints Over Multiple Relations

CREATE TABLE Sailors

(sid INTEGER,
sname CHAR(10),
rating INTEGER,
age REAL,
PRIMARY KEY (sid),
CHECK

((SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid) FROM
Boats B) < 100)

*Number of boats
plus number of
sailors is < 100*

- Awkward and wrong!
- Only checks sailors!
- Only required to hold if the associated table is non-empty.

- ASSERTION is the right solution; not associated with either table.

- Unfortunately, not supported in many DBMS.

- Triggers are another solution.

CREATE ASSERTION smallClub

CHECK

((SELECT COUNT (S.sid) FROM Sailors S)
+ (SELECT COUNT (B.bid)
FROM Boats B) < 100)



Or, Use a Trigger

- Trigger: procedure that starts automatically if specified changes occur to the DBMS
- Three parts:
 - Event (activates the trigger)
 - Condition (tests whether the triggers should run)
 - Action (what happens if the trigger runs)
- Triggers (in some form) are supported by most DBMSs; Assertions are not.
- Support for triggers is defined in the SQL:1999 standard.

Triggers

```
CREATE TRIGGER trigger_name
ON TABLE
{FOR {[INSERT][,][UPDATE][,][DELETE]}
[WITH APPEND]
AS
sql-statements
```

- Cannot be called directly – initiated by events on the database.
- Can be *synchronous* or *asynchronous* with respect to the transaction that causes it to be fired.

Triggers: Example

```
CREATE TRIGGER member_delete
ON member FOR DELETE
AS
IF (Select COUNT (*) FROM loan INNER JOIN deleted
    ON loan.member_no = deleted.member_no) > 0
BEGIN
    PRINT 'ERROR - member has books on loan.'
    ROLLBACK TRANSACTION
END
ELSE
DELETE reservation WHERE reservation.member_no
    = deleted.member_no
```

Summary: Triggers, Assertions, Constraints

- Very vendor-specific (although standard has been developed).
- Triggers vs. Constraints and Assertions:
 - Triggers are “operational”, others are declarative.
- Triggers can make the system hard to understand if not used with caution.
 - ordering of multiple triggers
 - recursive/chain triggers
- Triggers can be hard to optimize.
- But, triggers are also very powerful.
- Use to create high-performance, “*active*” databases.

Summary of Conceptual Design

- ***Conceptual design follows requirements analysis,***
 - Yields a high-level description of data to be stored
- **ER model popular for conceptual design**
 - Constructs are expressive, close to the way people think about their applications.
 - Note: There are many variations on ER model
 - Both graphically and conceptually
- **Basic constructs: *entities, relationships, and attributes* (of entities and relationships).**
- **Some additional constructs: *weak entities, ISA hierarchies, and aggregation.***

- **Several kinds of integrity constraints:**
 - *key constraints*
 - *participation constraints*
 - *overlap/covering* for ISA hierarchies.
- Some **foreign key constraints** are also implicit in the definition of a relationship set.
- Many other constraints (notably, *functional dependencies*) cannot be expressed.
- Constraints play an important role in determining the best database design for an enterprise.

Summary of ER (Cont.)

- ER design is *subjective*. There are often many ways to model a given scenario!
- Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
 - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, aggregation.
- Ensuring good database design: resulting relational schema should be analyzed and refined further.
 - Functional Dependency information and normalization techniques are especially useful.